

CS11-711 Advanced NLP

# Debugging and Understanding NLP Models

Graham Neubig



**Carnegie Mellon University**

Language Technologies Institute

w/ Some Slides by Danish Pruthi

# A Typical Situation

- You've implemented an NLP system based on neural networks
- You've looked at the code, and it looks OK
- It has low accuracy, or makes incomprehensible errors
- **What do I do?**

# Three Model Understanding Dimensions

- **Debugging Implementation:** Identifying problems in your implementation (or assumptions)
- **Actionable Evaluation:** Identifying typical error cases and understanding how to fix them
- **Interpreting Predictions:** Examining individual predictions to dig deeper

# Debugging

# In Neural Net Models, Debugging is Paramount!

- Models are often **complicated and opaque**
- **Everything is a hyperparameter** (network size, model variations, batch size/strategy, optimizer/learning rate)
- Non-convex, stochastic optimization has **no guarantee of decreasing/converging loss**

# Possible Causes

- **Training time problems**
  - Lack of model capacity
  - Poor training algorithm
  - Training time bug
- **Test time problems**
  - Disconnect between training and test
  - Failure of search algorithm
- **Overfitting**
- **Mismatch between optimized function and eval**

Don't debug all at once! Start top and work down.

Debugging at Training Time

# Identifying Training Time Problems

- Look at the **loss function** calculated on the **training set**
  - Is the loss function going down?
  - Is it going down basically to zero if you run training long enough (e.g. 20-30 epochs)?
  - If not, does it go down to zero if you use very small datasets?




# Is My Model Too Weak?

- Larger models tend to perform better, esp. when pre-trained (e.g. Raffel et al. 2020)

Model	GLUE Average	CoLA Matthew's	SST-2 Accuracy	MRPC F1	MRPC Accuracy	STS-B Pearson	STS-B Spearman
Previous best	89.4 <sup>a</sup>	69.2 <sup>b</sup>	97.1 <sup>a</sup>	<b>93.6<sup>b</sup></b>	<b>91.5<sup>b</sup></b>	92.7 <sup>b</sup>	92.3 <sup>b</sup>
T5-Small	77.4	41.0	91.8	89.7	86.6	85.6	85.0
T5-Base	82.7	51.1	95.2	90.7	87.5	89.4	88.6
T5-Large	86.4	61.2	96.3	92.4	89.9	89.9	89.2
T5-3B	88.5	67.1	97.4	92.5	90.0	90.6	89.8
T5-11B	<b>90.3</b>	<b>71.6</b>	<b>97.5</b>	92.8	90.4	<b>93.1</b>	<b>92.8</b>

- Larger models can learn with fewer steps (Kaplan et al. 2020, Li et al. 2020)



# Trouble w/ Optimization

- If increasing model size doesn't help, you may have an optimization problem
- Check your
  - **optimizer** (Adam? standard SGD?)
  - **learning rate** (is the rate you're using standard, are you using decay?)
  - **initialization** (uniform? Glorot?)
  - **minibatching** (are you using sufficiently large batches?)
- Pay attention to these details when replicating previous work

Debugging at Test Time

# Training/Test Disconnects

- Usually your loss calculation and prediction will be implemented in different functions
- Especially true for structured prediction models (e.g. encoder-decoders)
- Like all software engineering: **duplicated code is a source of bugs!**
- Also, usually loss calculation is minibatched, generation not.

# Debugging Minibatching

- Debugging mini-batched loss calculation
  - Calculate loss with **large batch size** (e.g. 32)
  - Calculate loss for **each sentence individually and sum**
  - The values should be the same (modulo numerical precision)
- Create a unit test that tests this!

# Debugging Structured Generation

- Your decoding code should get the same score as loss calculation
- Test this:
  - Call **decoding function**, to generate an output, and keep track of its score
  - Call **loss function** on the generated output
  - The score of the two functions should be the same
- Create a unit test doing this!

# Debugging Search

- As you make search better, the **model score** should get better (almost all the time)
- Search w/ varying beam sizes and make sure you get a better overall model score with larger sizes
- Create a unit test testing this!

Mismatch b/t Optimized  
Function and Evaluation Metric




# Loss Function, Evaluation Metric

- It is very common to optimize for maximum likelihood for training
- But even though likelihood is getting better, accuracy can get worse

# Example w/ Classification

- Loss and accuracy are de-correlated (see dev)




- Why? Model gets more confident about its mistakes.

# Managing Loss Function/ Eval Metric Differences

- Most principled way: use structured prediction techniques to be discussed in future classes
  - Structured max-margin training
  - Minimum risk training
  - Reinforcement learning
  - Reward augmented maximum likelihood

# A Simple Method: Early Stopping w/ Eval Metric



# Actionable Evaluation

# Look At Your Data!

- Both bugs and research directions can be found by **looking at your model outputs**
- Your model is repeating all the time
  - > I thought it was bad bad bad bad bad bad bad
  - need a new inference algorithm?
- The model is consistently failing on named entities
  - need a better model of named entities?

# Which Data to Look At?

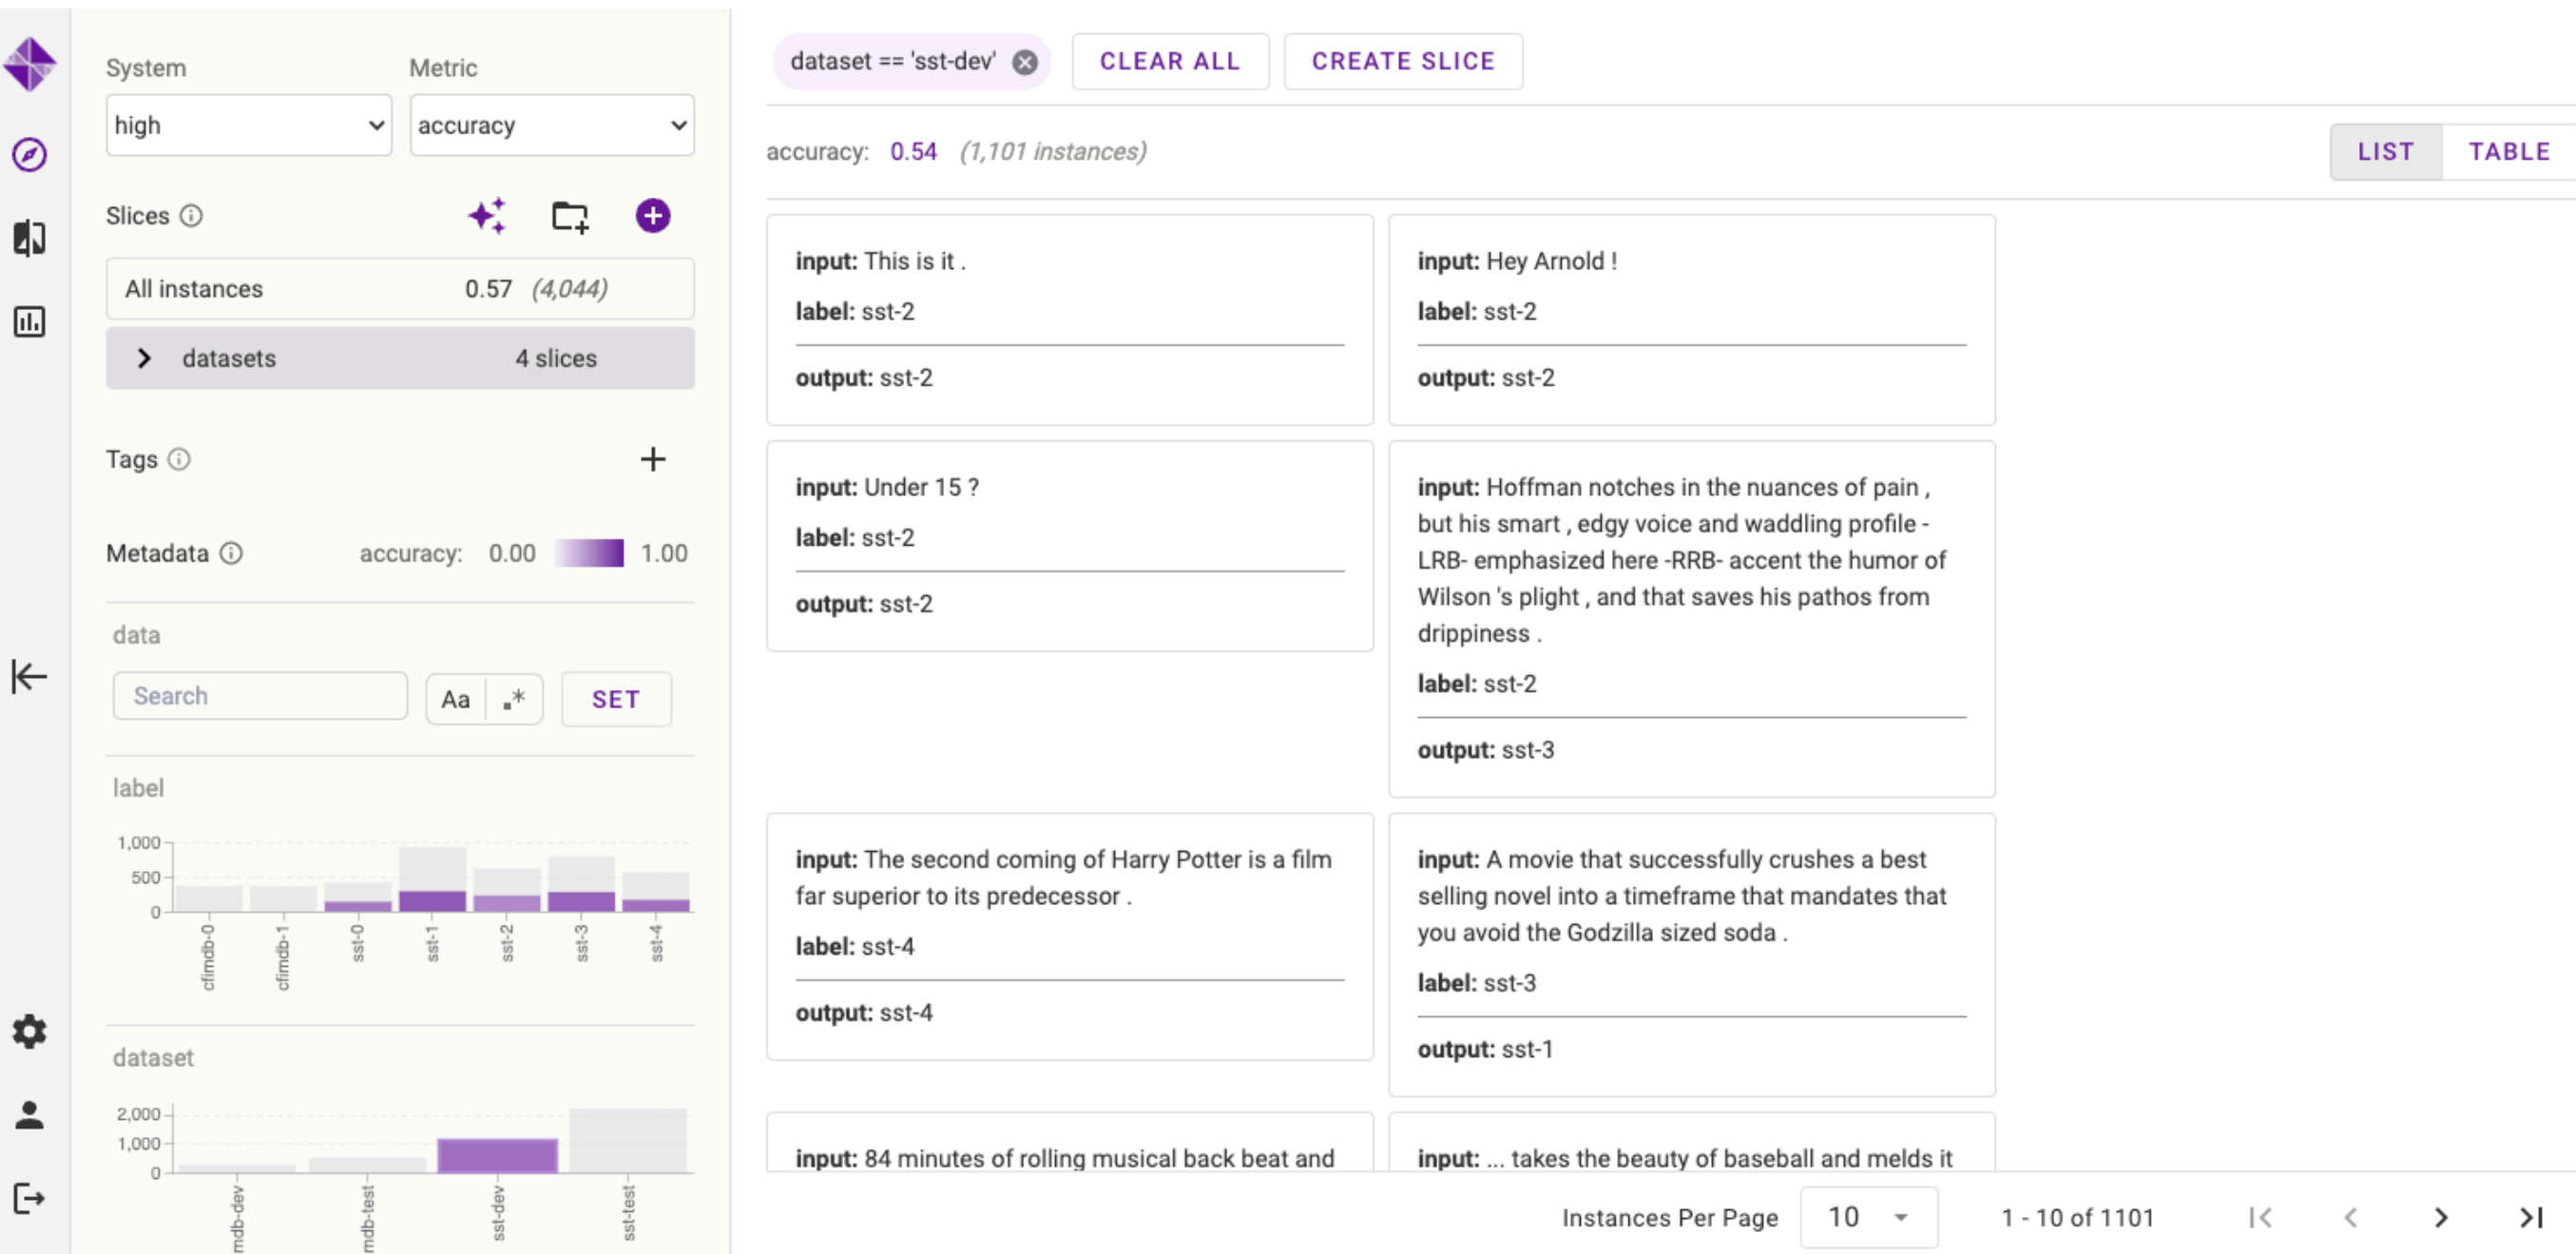
- Random examples
- Low-scoring examples ( $\text{low eval}(y)$ )
- Comparatively low examples ( $\text{low eval}(y_1) - \text{eval}(y_2)$ )

# Slicing

- Create a subset of your examples where you expect one model to do better than others
  - Long sentences
  - Sentences that contain a word
  - Sentences that belong to a cluster
  - etc. etc.



# Example: Zeno



<http://zenoml.com>

# Interpretation of Predictions and Model Internals

# Why Interpret Model Predictions?

- e.g. You want to know
  - **which words were used** in making a decision to verify its accuracy.
  - whether your model has learned a difficult pattern, or is focused on **spurious correlations**.
  - understand what information a **pre-trained model has captured** internally.

# LIME: Local Perturbations

For	Christmas	Song	visit	my	channel!	;)	prob	weight
1	0	1	1	0	0	1	0.17	0.57
0	1	1	1	1	0	1	0.17	0.71
1	0	0	1	1	1	1	0.99	0.71
1	0	1	1	1	1	1	0.99	0.86
0	1	1	1	0	0	1	0.17	0.57

label_prob	feature	feature_weight
0.9939024	channel!	6.180747
0.9939024	For	0.000000
0.9939024	;)	0.000000

# Explanation Technique: Gradient-based Scores

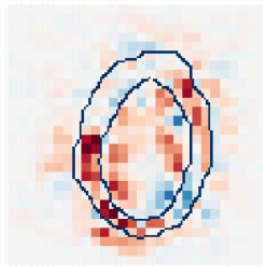
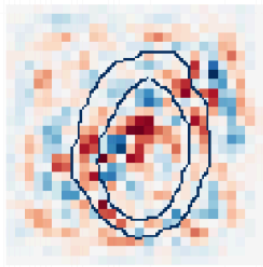
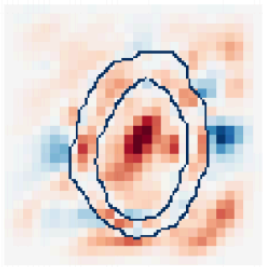
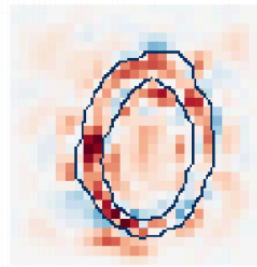
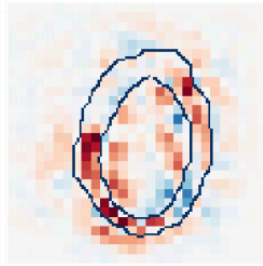
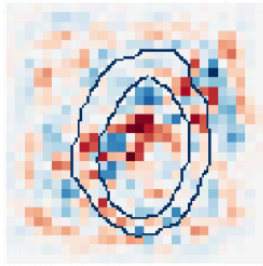
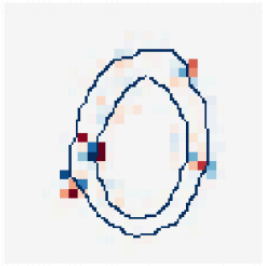
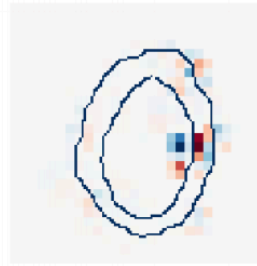
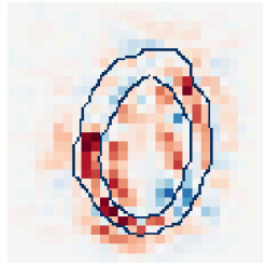
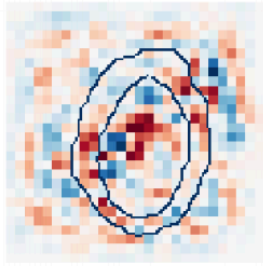
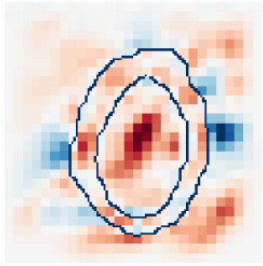
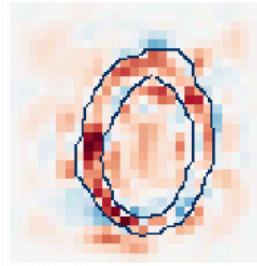

Method	Attribution $R_i^c(x)$	Example of attributions on MNIST			
Gradient * Input	$x_i \cdot \frac{\partial S_c(x)}{\partial x_i}$	ReLU	Tanh	Sigmoid	Softplus
Integrated Gradient	$(x_i - \bar{x}_i) \cdot \int_{\alpha=0}^1 \frac{\partial S_c(\tilde{x})}{\partial (\tilde{x}_i)} \bigg _{\tilde{x}=\bar{x}+\alpha(x-\bar{x})} d\alpha$				
<u><math>\epsilon</math>-LRP</u>	$x_i \cdot \frac{\partial^g S_c(x)}{\partial x_i}, \quad g = \frac{f(z)}{z}$				
<u>DeepLIFT</u>	$(x_i - \bar{x}_i) \cdot \frac{\partial^g S_c(x)}{\partial x_i}, \quad g = \frac{f(z) - f(\bar{z})}{z - \bar{z}}$				

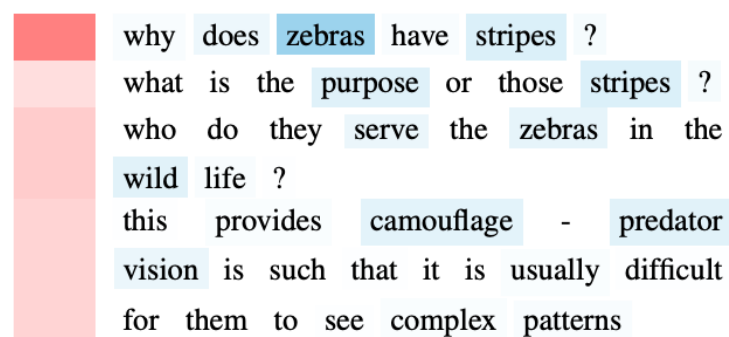
Figure from Ancona et al, ICLR 2018

# Explanation Technique: Attention



## Entailment

Rocktäschel et al, 2015

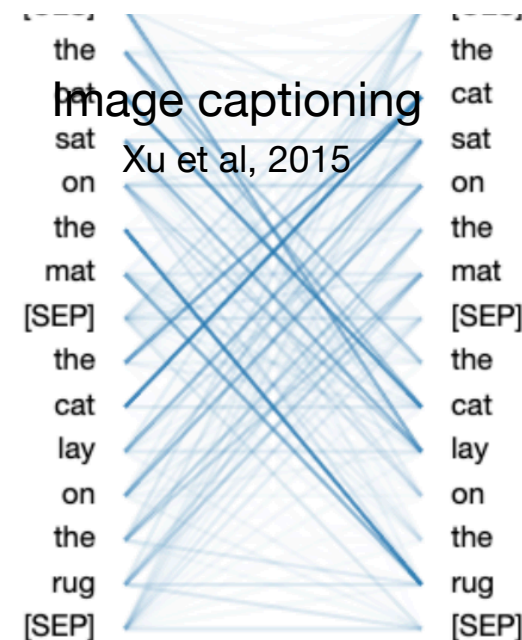


## Document classification

Yang et al, 2016



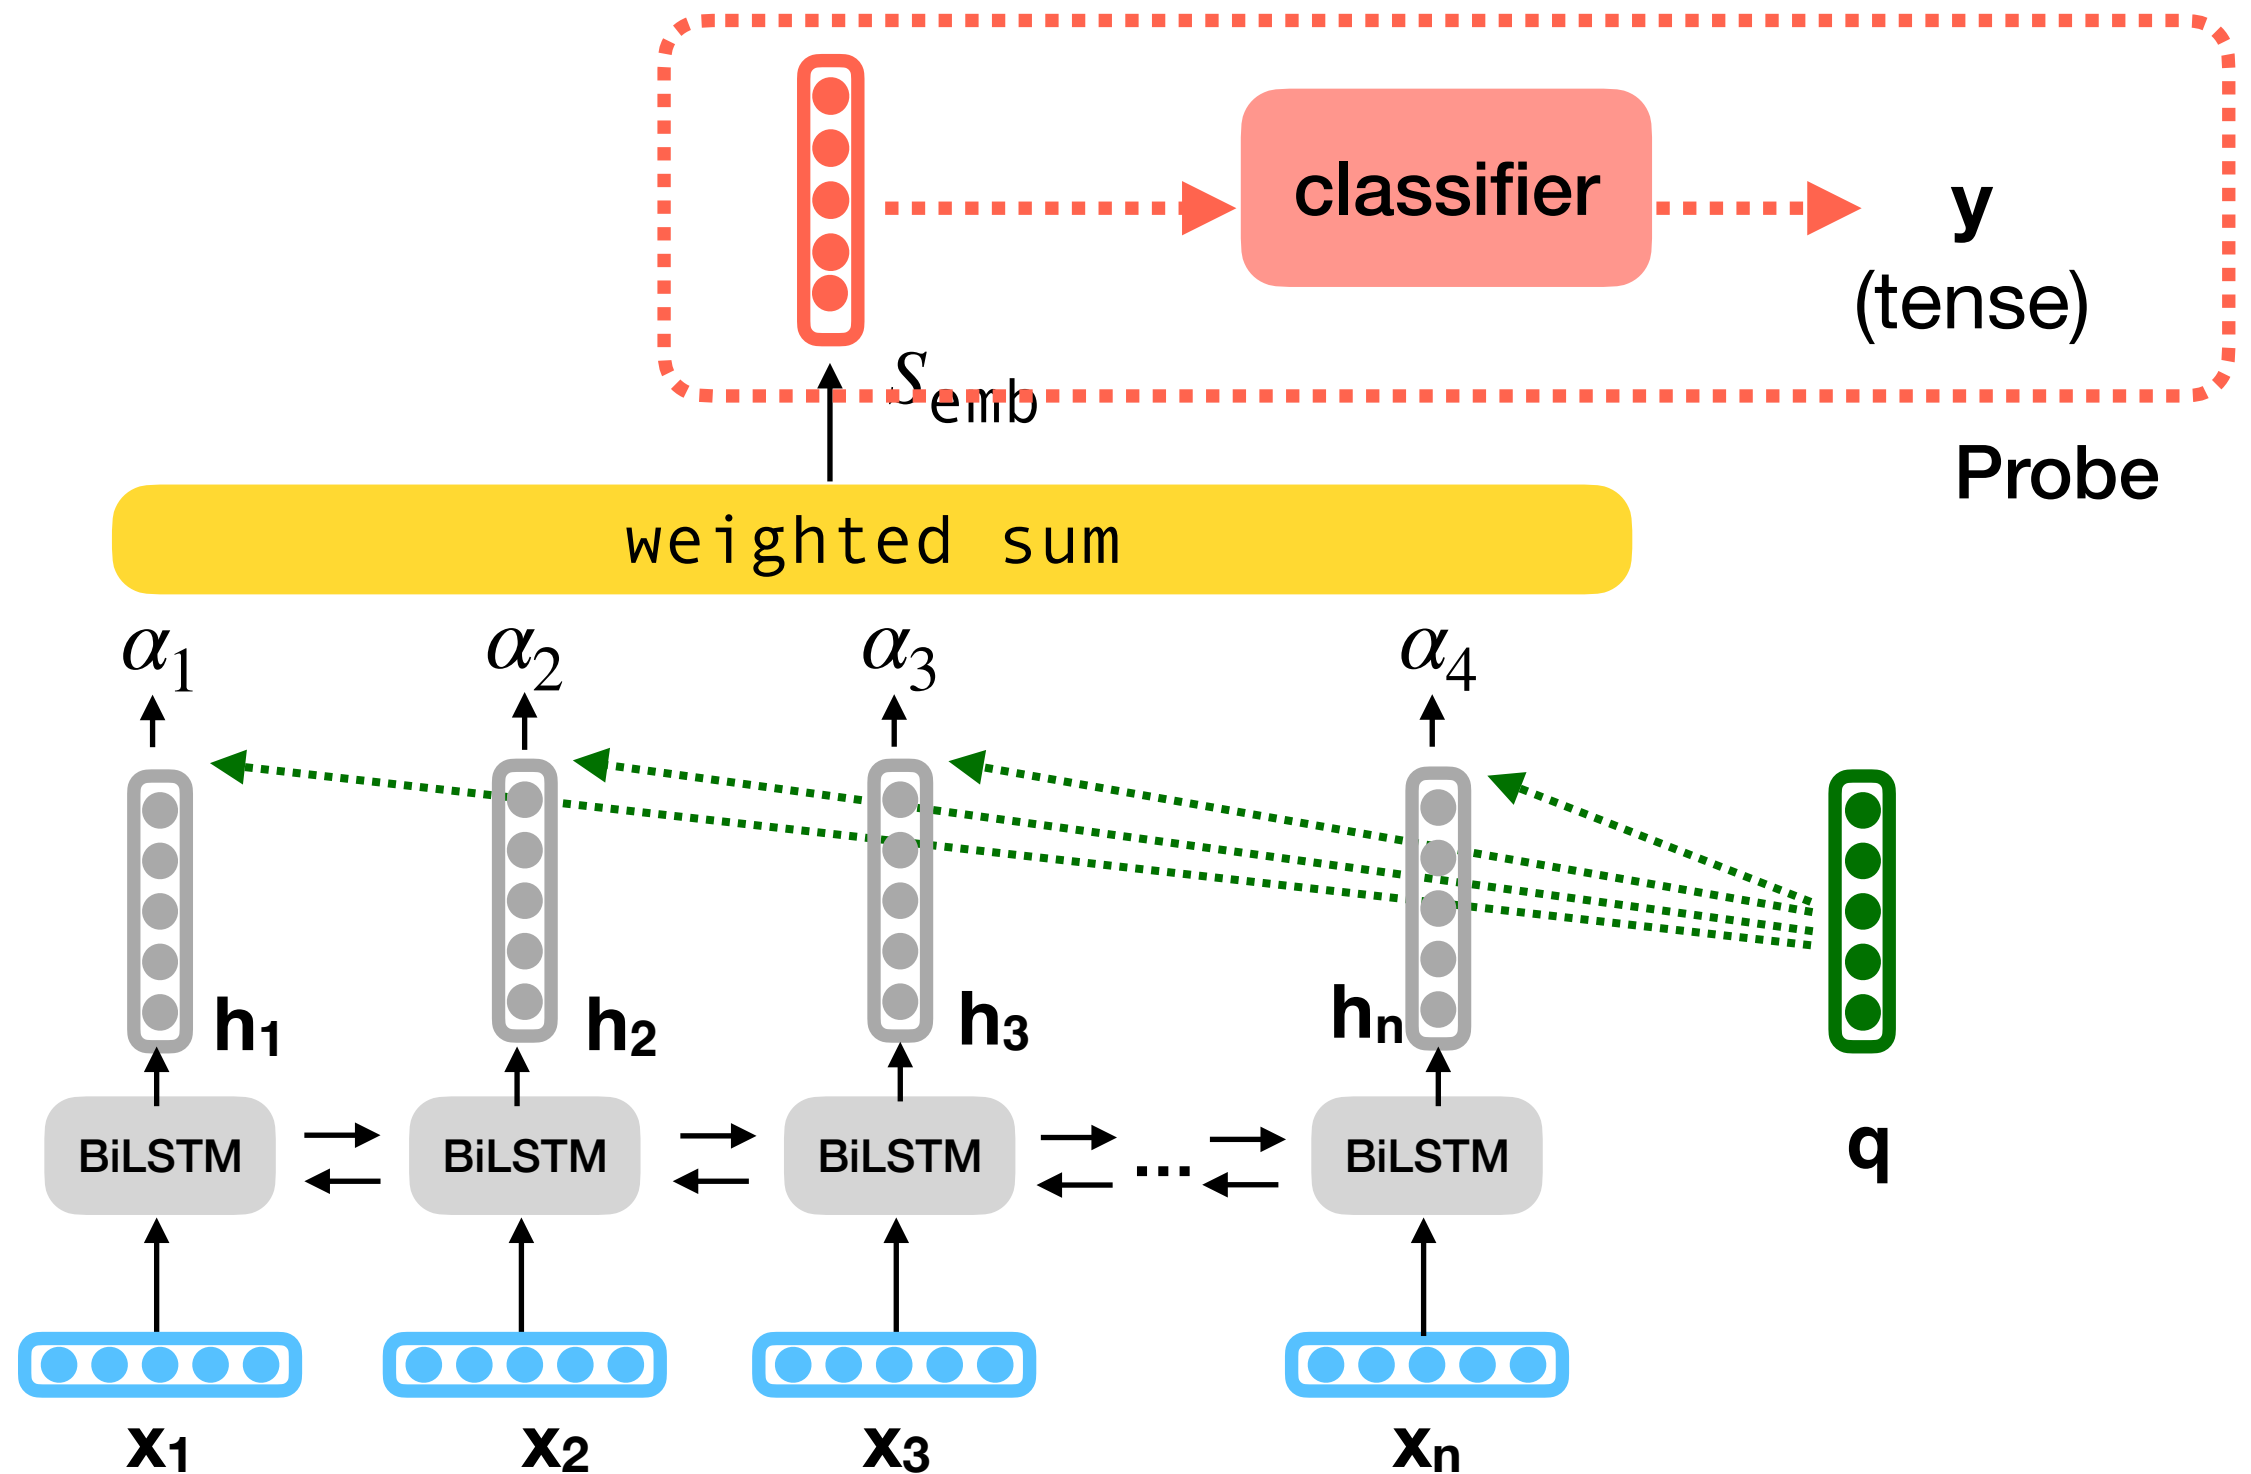
A stop sign is on a road with a mountain in the background.



## BERTViz

Vig et al, 2019

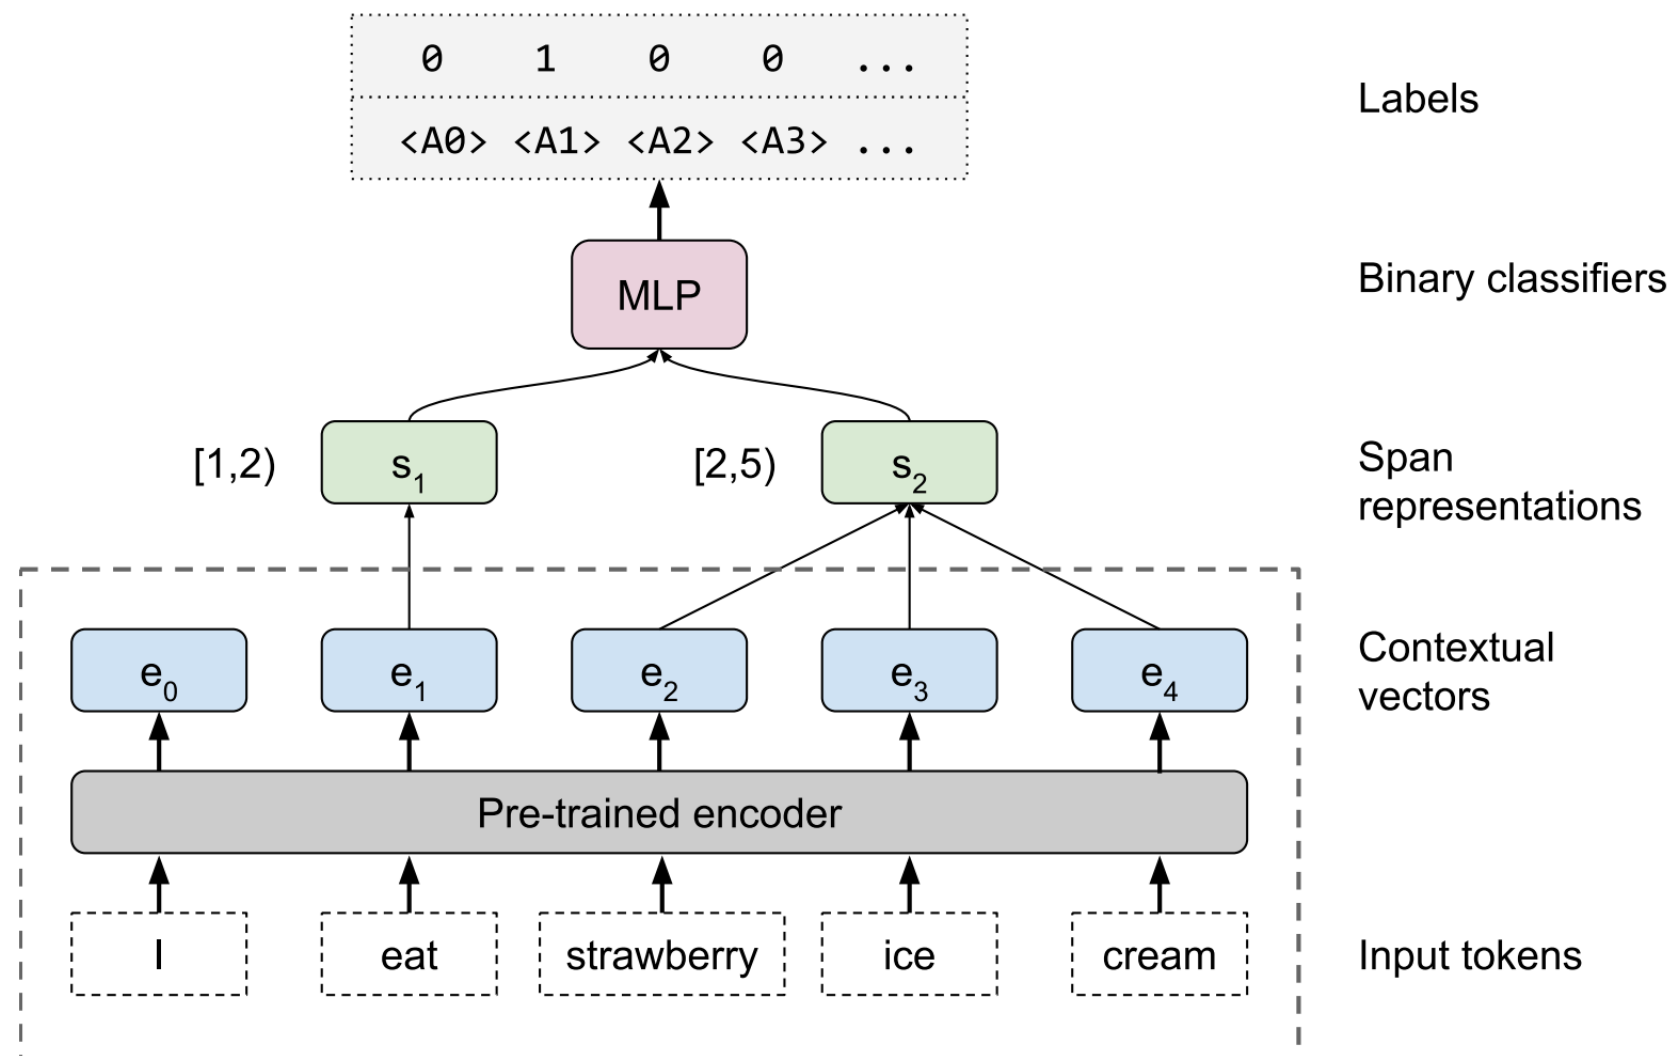
# Probing



# Edge Probing

(Tenney et al. 2019)

- A general framework that allows for probing of many types of information





# Issues with probing

- Did I interpret the representation or my probing classifier learn the task itself (Hewitt et al. 2019)
  - Solution - information theoretic probing that controls for classifier complexity (Voita et al. 2020)
- Can only probe for properties you have supervision for
- Correlation doesn't imply causation
- and more...

Questions?