


Natural language to Code Generation

Zora Wang, Nikitha Rao
Nov 9th, 2023

The NL2Code Task


Given a natural language instruction Q , generate code implementation C



The Landscape for NL2Code Generation

- Transition of Evaluation Metrics:

- Lexical, neural based metrics
- Test case execution



- Domain Coverage


- Built-in grammar: "sum([1, 2, 4])"
- Domain-specific: data science
- Open domain: diverse Python libraries

- Functional Complexity

- Simple (toy) functions: e.g., LeetCode
- Class level
- Repository level

- Test Automation

- Human-written tests
- Fuzzing methods
- Integrating LLMs



What We Had for Code Generation

Benchmark: **CodeXGLUE**, CoNaLa

- [code-code](#) (clone detection, defect detection, cloze test, code completion, code repair, and code-to-code translation)
- [text-code](#) (natural language code search, text-to-code generation)
- [code-text](#) (code summarization)
- [text-text](#) (documentation translation)

Model

- CodeBERT, GraphCodeBERT

Evaluation Metrics

- Lexical based: BLEU
- Syntax based: CodeBLEU

Category	Task	Dataset Name	Language	Train/Dev/Test Size	Baselines	Task definition
Code-Code	Clone Detection	BigCloneBench	Java	900K/416K/416K	CodeBERT	Predict semantic equivalence for a pair of codes.
		POJ-104	C/C++	32K/8K/12K		Retrieve semantically similar codes.
	Defect Detection	Devign	C	21k/2.7k/2.7k		Identify whether a function is vulnerable.
		CT-all	Python, Java, PHP, JavaScript, Ruby, Go	-/-/176k		Tokens to be predicted come from the entire vocab.
	Cloze Test	CT-max/min	Python, Java, PHP, JavaScript, Ruby, Go	-/-/2.6k	CodeGPT	Tokens to be predicted come from {max, min}.
		PY150	Python	100k/5k/50k		
	Code Completion	GitHub Java Corpus	Java	13k/7k/8k	CodeGPT	Predict following tokens given contexts of codes.
	Code Repair	Bugs2Fix	Java	98K/12K/12K	Encoder-Decoder	Automatically refine codes by fixing bugs.
Text-Code	NL Code Search	CodeSearchNet, AdvTest	Python	251K/9.6K/19K	CodeBERT	Given a natural language query as input, find semantically similar codes.
		CodeSearchNet, WebQueryTest	Python	251K/9.6K/1k		Given a pair of natural language and code, predict whether they are relevant or not.
	Text-to-Code Generation	CONCODE	Java	100K/2K/2K	CodeGPT	Given a natural language docstring/comment as input, generate a code.
Code-Text	Code Summarization	CodeSearchNet	Python, Java, PHP, JavaScript, Ruby, Go	908K/45K/53K	Encoder-Decoder	Given a code, generate its natural language docstring/comment.
Text-Text	Documentation Translation	Microsoft Docs	English-Latvian/Danish/Norwegian/Chinese	156K/4K/4K		Translate code documentation between human languages (e.g. En-Zh), intended to test low-resource multi-lingual translation.

What We Had for Code Generation

Benchmark:


- CodeXGLUE, CoNaLa

Model



- CodeBERT, GraphCodeBERT

Evaluation Metrics

- Lexical based: BLEU
- Syntax based: **CodeBLEU**



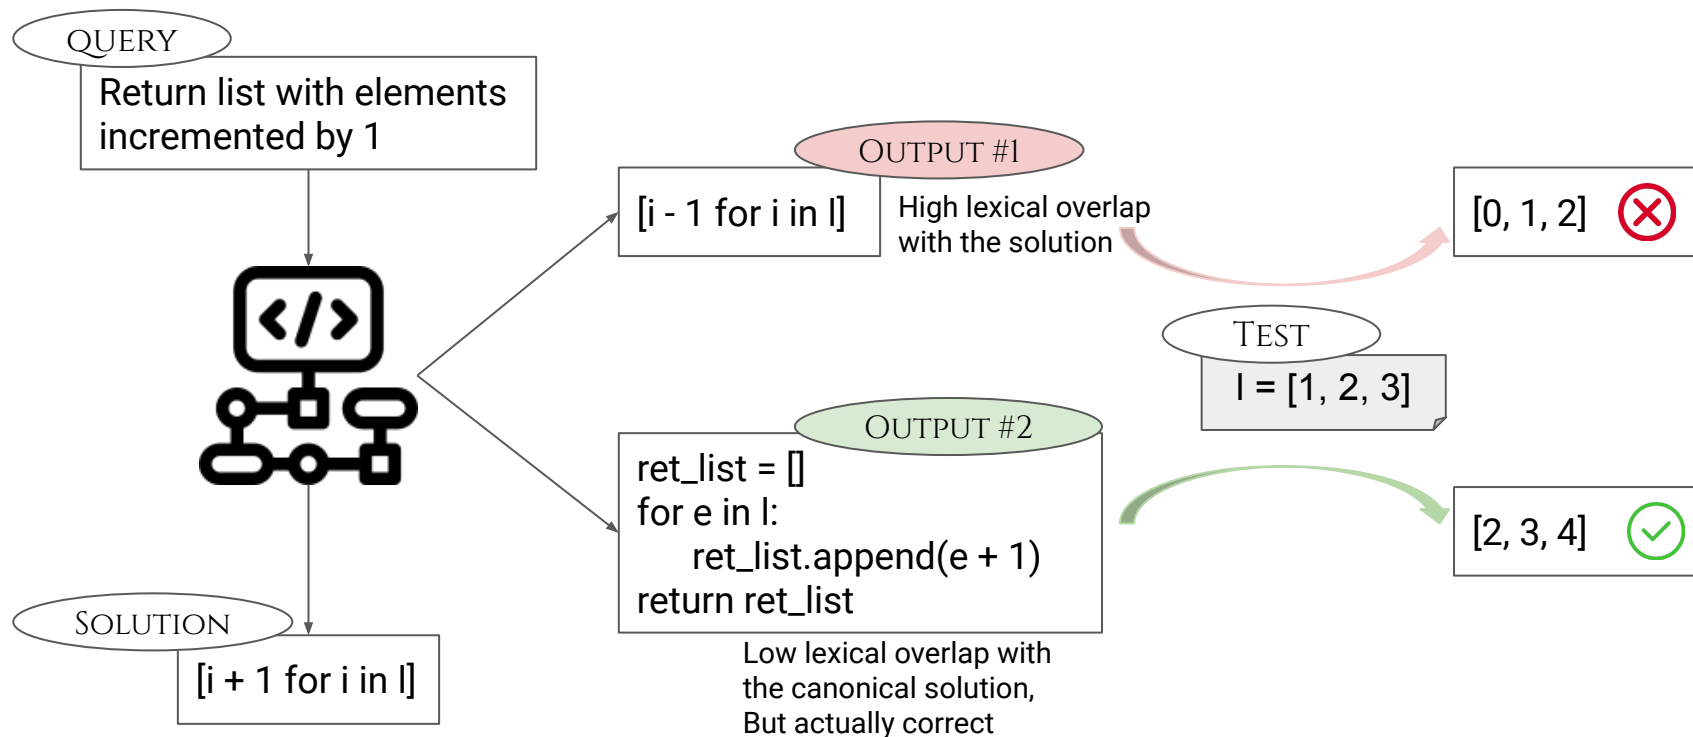
Weighted N-Gram Match



Semantic Data-flow Match

$$\text{CodeBLEU} = \alpha \cdot \text{N-Gram Match (BLEU)} + \beta \cdot \text{Weighted N-Gram Match} + \gamma \cdot \text{Syntactic AST Match} + \delta \cdot \text{Semantic Data-flow Match}$$

Issues: Evaluations Are Not Rigorous

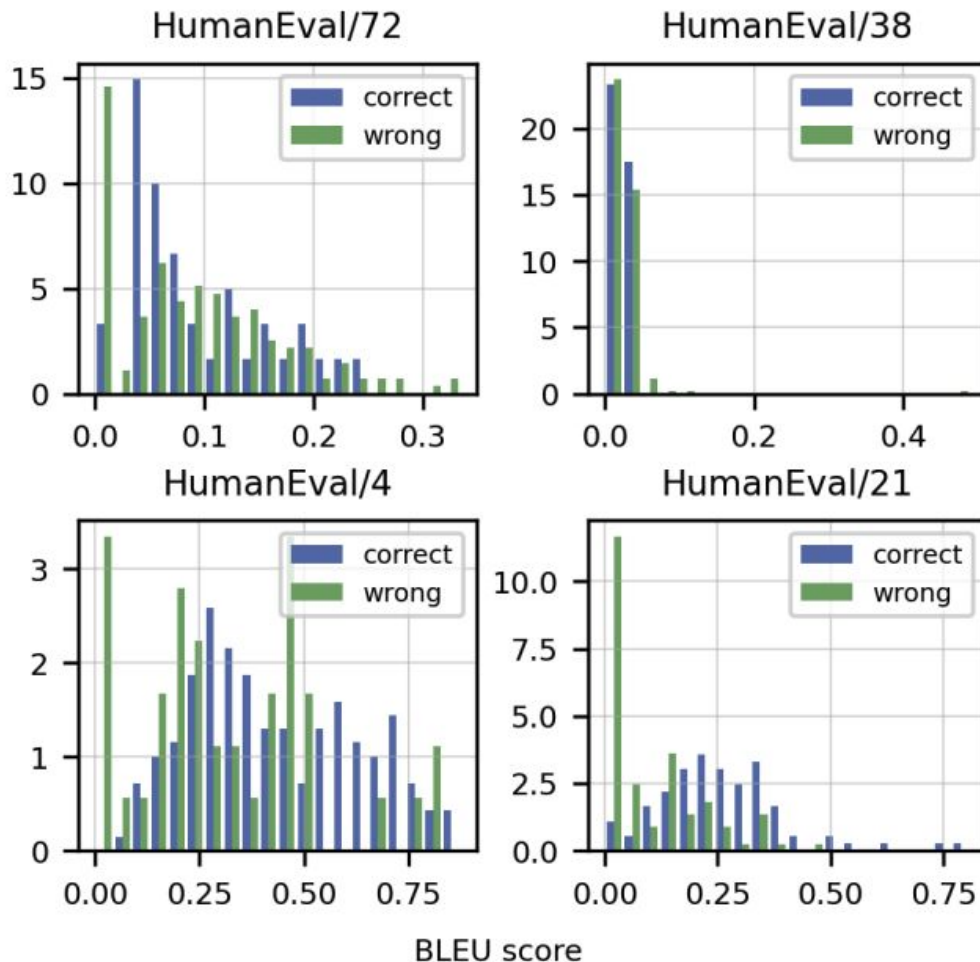


HumanEval Benchmark -

Evaluation: test case execution

164 hand-written examples

- Why human-written?
 - “It is important for these tasks to have a small fraction of GitHub, which already contains a large number of programs that are correct and safe to execute.”
- Safety: Sandbox for Executing Programs
 - “Since publicly available programs are often incorrect, executing these programs in a sandbox is necessary to avoid damage to the system.”
- Optimizing BLEU \neq Improving Program Quality



HumanEval Looks Like Toy Examples?

HumanEval Examples

```
def incr_list(l: list):
    """Return list with elements incremented by 1.
    >>> incr_list([1, 2, 3])
    [2, 3, 4]
    >>> incr_list([5, 3, 5, 2, 3, 3, 9, 0, 123])
    [6, 4, 6, 3, 4, 4, 10, 1, 124]
    """
    return [i + 1 for i in l]
```

```
def solution(lst):
    """Given a non-empty list of integers, return the sum of all of the odd elements
    that are in even positions.

    Examples
    solution([5, 8, 7, 1]) ==>12
    solution([3, 3, 3, 3, 3]) ==>9
    solution([30, 13, 24, 321]) ==>20
    """
    return sum(lst[i] for i in range(0, len(lst)) if i % 2 == 0 and lst[i] % 2 == 1)
```

Real-World Development Code

Asking the user for input until they give a valid response

Asked 9 years, 6 months ago Modified 1 year, 5 months ago Viewed 1.0m times

I am writing a program that accepts user input.

750

```
#note: Python 2.7 users should use `raw_input`, the equivalent of 3.X's `input`
age = int(input("Please enter your age: "))
if age >= 18:
    print("You are able to vote in the United States!")
else:
    print("You are not able to vote in the United States.")
```



transformers Public

Watch 1.1k

Fork 22.9k

Starred 114k

main

262 branches 139 tags

Go to file

Add file

Code

About



hi-sushanta Removed the redundant SiLUActivation class. (#27136) ✓ 4991216 1 hour ago 14,388 commits

.circleci	Limit to inferior fsspec version (#27010)	last week
.github	Dev version	2 hours ago
docker	[core / Quantization] AWQ integration (#27045)	2 days ago
docs	translate peft.md to chinese (#27215)	2 hours ago
examples	Dev version	2 hours ago
model_cards	Update URL for Hub PR docs (#17532)	last year
notebooks	Update README.md (#25941)	2 months ago
scripts	Fix stale bot for linked issues (#26711)	3 weeks ago

Transformers: State-of-the-art Machine Learning for Pytorch, TensorFlow, and JAX.

huggingface.co/transformers

[python](#) [nlp](#) [machine-learning](#)
[natural-language-processing](#)
[deep-learning](#) [tensorflow](#) [pytorch](#)
[transformer](#) [speech-recognition](#)
[seq2seq](#) [flax](#) [pretrained-models](#)
[language-models](#) [nlp-library](#)
[language-model](#) [hacktoberfest](#) [bert](#)
[jax](#) [pytorch-transformers](#) [model-hub](#)

Domain Exploration

- Leetcode Style: HumanEval, APPS, MBPP
 - Manually written or collected from code contest websites
 - Only uses Python built-in grammar
- Limited Domains: e.g., Data Science
 - DS-1000: StackOverflow questions
 - ARCADE: Interactive Jupyter Notebooks
 - ExeDS
 -
- Open Domain: ODEX
 - 79 Python libraries
 - Four natural languages

NumPy SciPy Pandas TensorFlow PyTorch Scikit-learn Matplotlib

Problem:
I could not find a built-in function in Python to generate a log uniform distribution given a min and max value (the R equivalent is here), something like: `logunif(n, min, max, base)` that returns `n` log uniformly distributed in the range `min` and `max`.
The closest I found though was `numpy.random.uniform`.
That is, given range of `x`, I want to get samples of given size (`n`) that suit log-uniform distribution.
Any help would be appreciated!

A:
<code>
import numpy as np
min = 1
max = np.e
n = 10000
</code>
BEGIN SOLUTION
<code>
[insert]
</code>
END SOLUTION
<code>
print(result)
</code>

Reference Solution

```
import scipy.stats  
result = scipy.stats.loguniform.rvs(a = min, b = max, size = n)
```

Automatic Evaluation

Test case 1

```
min = 1  
max = np.e  
n = 10000  
ans = ... # generated by Reference solution
```

Test code

```
np.testing.assert_array_equal(result.shape, ans.shape)  
from scipy.stats import ks_2samp  
# Kolmogorov-Smirnov Test judges whether the two sampled  
# from similar distribution  
assert ks_2samp(result, ans)[0] <= 0.1
```

Surface form constraints
for and while should not appear in Syntax Tree

NumPy example problem involving randomness, requiring the use of a specialist knowledge test.

Execution-Based Evaluation for Open-Domain Code Generation

- Larger Domain Coverage



- Test execution on real-world coding queries

- Collected from StackOverflow questions

- Support four natural languages as input

- English, Spanish, Japanese, Russian

```
import requests
```

```
def function(files, url, data):
    """multipartのリクエストで複数のデータ`files`, `data`を`url`にPOSTする
    (POST multiple data `files`, `data` to `url` with multipart request)
    """
```

```
return requests.post(url, files=files, data=data)
```

```
# test case
r = requests.Response()
r.status_code = 200
requests.post = Mock(return_value = r)
file_path = 'a.txt'
```

Dataset	Samples	Domain	Executable?	Avg. Test Cases	Data Source	NL
JuICE (Agashe et al., 2019)	1,981	open	✗	-	GitHub Notebooks	en
HumanEval (Chen et al., 2021)	164	4	✓	7.7	Hand-written	en
MBPP (Austin et al., 2021)	974	8	✓	3.0	Hand-written	en
APPS (Hendrycks et al., 2021)	10,000	0	✓	13.2	Competitions	en
DSP (Chandel et al., 2022)	1,119	16	✓	2.1	Github Notebooks	en
MTPB (Nijkamp et al., 2022)	115	8	✓	5.0	Hand-written	en
Exe-DS (Huang et al., 2022)	534	28	✓	-	GitHub Notebooks	en
DS-1000 (Lai et al., 2022)	1,000	7	✓	1.6	StackOverflow	en
CoNaLa (Yin et al., 2018)	2,879	open	✗	-	StackOverflow	en
MCoNaLa (Wang et al., 2022)	896	open	✗	-	StackOverflow	es, ja, ru
ODEX	945	79	✓	1.8	StackOverflow Hand-Written	en, es, ja, ru

```
(calculate the improper integral given by the function f
from the number `n` to infinity)
"""
```

```
return
```

```
return sympy.integrate(f, (sympy.symbols('x'), n, sympy.oo))
```

```
# test case
x = sympy.symbols('x')
f = (x * x)
n = 1
assert str(function(f, n)) == 'oo'
```


ODEX: Unique Challenges for Execution

Closed-domain code: easy to execute and verify

```
assert func([1, 2, 10]) == [2, 3, 11]
```

Open-domain code:

- Random outputs



- Specialized verification

```
In [1]: import numpy as np
```

```
In [2]: a = np.array([1, 2, 3])
```

```
In [3]: b = np.array([1, 2, 3])
```

```
In [4]: assert (a == b)
```

ValueError

Cell In[4], line 1

----> 1 assert (a == b)

Traceback (most recent call last)

ValueError: The truth value of an array with more than one element is ambiguous.

```
In [5]: np.array_equal(a, b)
```

```
Out[5]: True
```

- (Potentially) not reproducible queries
 - HTTP requests, e.g., requests.post(["https://def.xyz"](https://def.xyz), data={'key': 'value'})

Significant Performance Gaps: Open vs. Closed

- Although Codex performs better overall
- CodeGen has smaller domain gaps




Figure 7: CODEX (left) and CODEGEN (right) *pass@1* on open- and closed-domain problems in each language.

Functional Complexity

- Function Level: HumanEval, MBPP
- Class Level: ClassEval

HumanEval Function Test

```
METADATA = {\n  'author': 'jt',\n  'dataset': 'test'\n}

def check(candidate):
    assert candidate([1.0, 2.0, 3.9, 4.0, 5.0, 2.2], 0.3) == True
    assert candidate([1.0, 2.0, 3.9, 4.0, 5.0, 2.2], 0.05) == False
    ...
```

MBPP Function Test

```
[
    "assert get_ludic(10) == [1, 2, 3, 5, 7]",
    "assert get_ludic(25) == [1, 2, 3, 5, 7, 11, 13, 17, 23, 25]", ...
]
```

ClassEval Method Test

```
class VendingMachineTestPurchaseItem(unittest.TestCase):
    def test_purchase_item(self):
        vm = VendingMachine()
        vm.inventory = {'Coke': {'price': 1.25, 'quantity': 10}}
        vm.balance = 1.25
        self.assertEqual(vm.purchase_item('Coke'), 0.0)
        self.assertEqual(vm.inventory, {'Coke': {'price': 1.25, 'quantity': 9}})

    def test_purchase_item_2(self):
        vm = VendingMachine()
        vm.inventory = {'Coke': {'price': 1.25, 'quantity': 10}}
        vm.balance = 1.25
        self.assertEqual(vm.purchase_item('Pizza'), False)
        self.assertEqual(vm.inventory, {'Coke': {'price': 1.25, 'quantity': 10}})
    ...
```

ClassEval Class Test

```
class VendingMachineTestMain (unittest.TestCase):
    def setUp(self) -> None:
        self.vm = VendingMachine()
        self.vm.inventory = {'Coke': {'price': 1.25, 'quantity': 10}}
        self.vm.balance = 0

    def test_all(self):
        self.assertEqual(vm.insert_coin(1.25), 1.25)
        self.assertEqual(vm.purchase_item('Coke'), 0.0)
        self.assertEqual(vm.inventory, {'Coke': {'price': 1.25, 'quantity': 9}})
        self.assertEqual(vm.restock_item('Coke', 10), True)
        self.assertEqual(vm.inventory, {'Coke': {'price': 1.25, 'quantity': 19}})
        self.assertEqual(vm.display_items(), 'Coke - $1.25 [19]')
    ...
```

```
from datetime import datetime
class VendingMachine:
    """This is a class to simulate a vending machine, including adding products, inserting
    coins, purchasing products, viewing balance, replenishing product inventory, and
    displaying product information. """
    def __init__(self):
        """
        Initializes the vending machine's inventory and balance.
        """
        self.inventory= []
        self.balance= {}
    def purchase_item(self, item_name):
        """ Purchases a product from the vending machine and returns the balance after the
        purchase.
```

Import Statements

Class Name

Class Description

Class Constructor

Method Signature

Functional Description

Test Fixtures:
setUp

Figure 4: Test Cases in Existing Benchmarks and ClassEval

"Write a python function to find the first repeated character in a given string."

Figure 1: Examples in Existing Benchmarks

```
returns false.
>>> vendingMachine.inventory = {'Coke': {'price': 1.25, 'quantity': 10}}
>>> vendingMachine.restock_item('Coke', 10)
True
>>> vendingMachine.inventory
{'Coke': {'price': 1.25, 'quantity': 20}}
"""
...
```

Parameter/Return Description

Example Input/Output

Figure 2: An Example of Class Skeleton in ClassEval

Functional Complexity

- Function Level: HumanEval, MBPP
- Class Level: ClassEval
- Repository Level:
 - RepoCoder
 - Retrieval-augmented generation
 - Multiple iterations
 - RepoEval
 - Collected 14 Github Repositories
 - Metrics:
 - exact match
 - exact similarity
 - execution

```
# Below are some referential code fragments      Retrieved Code
from other files:
# -----
# the below code fragment can be found in:
# tests/test_pipelines_common.py
# -----
# @unittest.skipIf(torch_device != "cuda")
# def test_to_device(self):
#     components = self.get_dummy_components()
#     pipe = self.pipeline_class(**components)
#     pipe.progress_bar(disable=None)
#     pipe.to("cpu")
# -----
"""Based on above, complete the following code:"""

@unittest.skipIf(torch_device != "cuda")      Unfinished Code
def test_float16_inference(self):
    components = self.get_dummy_components()

    pipe = self.pipeline_class(**components)    Model
    pipe.to(torch_device)                       Prediction
```

Figure 3: A visual example demonstrating the format of the RepoCoder prompt, which combines the retrieved code snippets from the repository with the unfinished code present in the target file.

Functional Complexity

- Function Level: HumanEval, MBPP
- Class Level: ClassEval
- Repository Level: RepoCoder & RepoEval
- Repo-Level Pull Requests: SWE-Bench

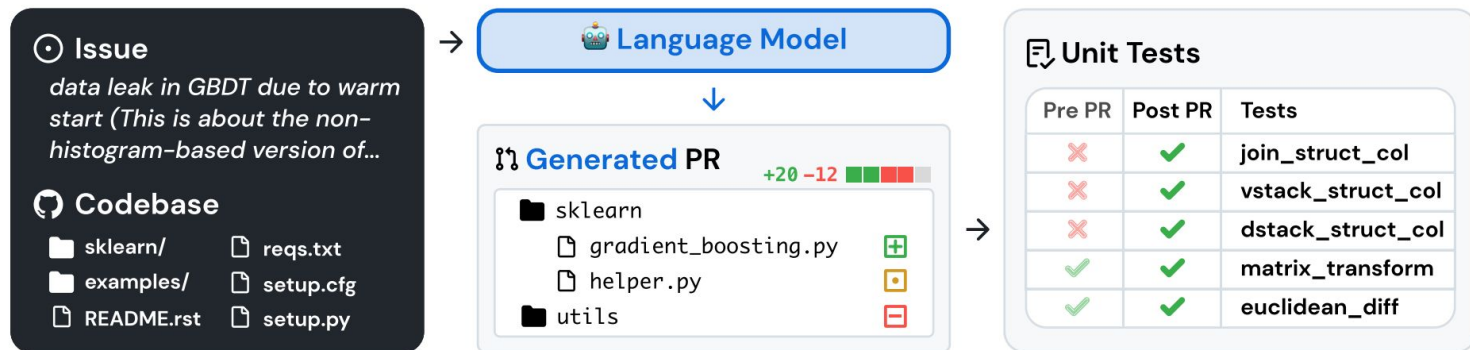


Figure 1: SWE-bench sources task instances from real-world Python repositories by connecting GitHub issues to merged pull request solutions that resolve related tests. Provided with the issue text and a codebase snapshot, models generate a patch that is evaluated against real tests.

Model Input

▼ Instructions

• 1 line

You will be provided with a partial code base and an issue statement explaining a problem to resolve.

▼ Issue

• 67 lines

napoleon_use_param should also affect "other parameters" section Subject: napoleon_use_param should also affect "other parameters" section

Problem

Currently, napoleon always renders the Other parameters section as if napoleon_use_param was False, see source

```
def _parse_other_parameters_section(self, section: str) -> List[str]:
    # type: (unicode) -> List[unicode]
    return self._format_fields(_('Other Parameters'), self._consume_fields())

def _parse_parameters_section(self, section: str) -> List[str]:
    # type: (unicode) -> List[unicode]
    fields = self._consume_fields()
    if self._config.napoleon_use_param: ...
```

▼ Code

• 1431 lines

► README.rst

• 132 lines

► sphinx/ext/napoleon/docstring.py

• 1295 lines

► Additional Instructions

• 57 lines

Gold Patch

sphinx/ext/napoleon/docstring.py

```
def _parse_other_parameters_section(self, section: str) -> List[str]:
- return self._format_fields(_('Other Parameters'), self._consume_fields())
+ if self._config.napoleon_use_param:
+     # Allow to declare multiple parameters at once (ex: x, y: int)
+     fields = self._consume_fields(multiple=True)
+     return self._format_docutils_params(fields)
+ else:
+     fields = self._consume_fields()
+     return self._format_fields(_('Other Parameters'), fields)
```

Generated Patch

sphinx/ext/napoleon/docstring.py

```
def _parse_other_parameters_section(self, section: str) -> List[str]:
- return self._format_fields(_('Other Parameters'), self._consume_fields())
+ return self._format_docutils_params(self._consume_fields())
```

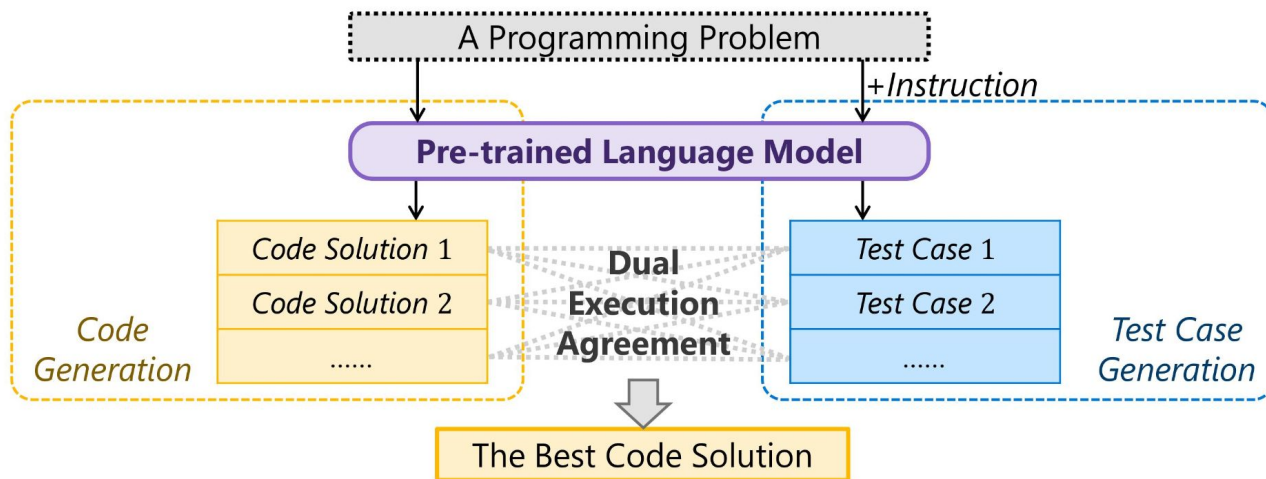
Generated Patch Test Results

```
PASSED NumpyDocstringTest (test_yield_types)
PASSED TestNumpyDocstring (test_escape_args_and_kwargs 1)
PASSED TestNumpyDocstring (test_escape_args_and_kwargs 2)
PASSED TestNumpyDocstring (test_escape_args_and_kwargs 3)
PASSED TestNumpyDocstring (test_pep526_annotations)
FAILED NumpyDocstringTest (test_parameters_with_class_reference)
FAILED TestNumpyDocstring (test_token_type_invalid)
===== 2 failed, 45 passed, 8 warnings in 5.16s =====
```

Figure 6: We show an example of a formatted task instance, a model prediction, and the testing framework logs. Results and inputs are stylized for readability. In the gold and generated patch file, red-highlighted lines represent deletions and green-highlighted lines represent additions.

Automated & Improved Testing

- Initial approach: human (expert) write test cases
- Evaluate on code without annotated tests? CodeT
 - Models generate solutions and test cases at the same time
 - Cross-validation between multiple solutions and tests



Automated & Improved Testing

- Initial approach: human (expert) write test cases
- LLMs to create test cases → EvalPlus
 - Scale-up test generation, reduce human effort
 - More sufficient tests, can find wrong solutions

Table 2: Overview of EvalPlus-improved benchmarks.

	#Tests				#Tasks
	Avg.	Medium	Min.	Max.	
HUMANEVAL	9.6	7.0	1	105 ²	
HUMANEVAL ⁺	764.1	982.5	12	1,100	164
HUMANEVAL ⁺ -MINI	16.1	13.0	5	110	

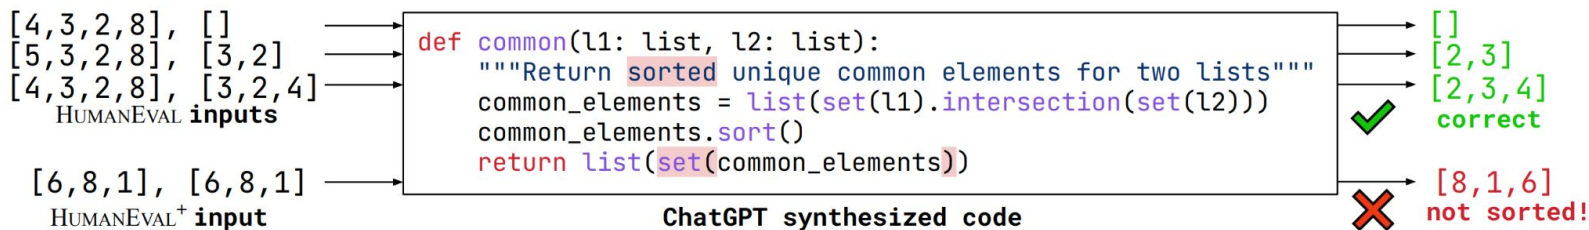


Figure 1: Exemplary wrong code synthesized by ChatGPT for HUMANEVAL #58.

Automated & Improved Testing

- Initial approach: human (expert) write test cases
- LLMs to create test cases → EvalPlus
- Fuzzing: a common method in software engineering
 - Type-aware value mutations / alterations
 - Limitation: can only apply to Python basic types, not open-domain ones, e.g., numpy array

Table 1: List of basic type-aware mutations over input x .

Type	Mutation	Type	Mutation
<code>int float</code>	Returns $x \pm 1$	<code>List</code>	$\left\{ \begin{array}{l} \text{Remove/repeat a random item } x[i] \\ \text{Insert/replace } x[i] \text{ with } \text{Mutate}(x[i]) \end{array} \right.$
<code>bool</code>	Returns a random boolean	<code>Tuple</code>	Returns <code>Tuple(Mutate(List(x)))</code>
<code>NoneType</code>	Returns <code>None</code>	<code>Set</code>	Returns <code>Set(Mutate(List(x)))</code>
<code>str</code>	$\left\{ \begin{array}{l} \text{Remove a sub-string } s \\ \text{Repeat a sub-string } s \\ \text{Replace } s \text{ with } \text{Mutate}(s) \end{array} \right.$	<code>Dict</code>	$\left\{ \begin{array}{l} \text{Remove a key-value pair } k \rightarrow v \\ \text{Update } k \rightarrow v \text{ to } k \rightarrow \text{Mutate}(v) \\ \text{Insert } \text{Mutate}(k) \rightarrow \text{Mutate}(v) \end{array} \right.$